



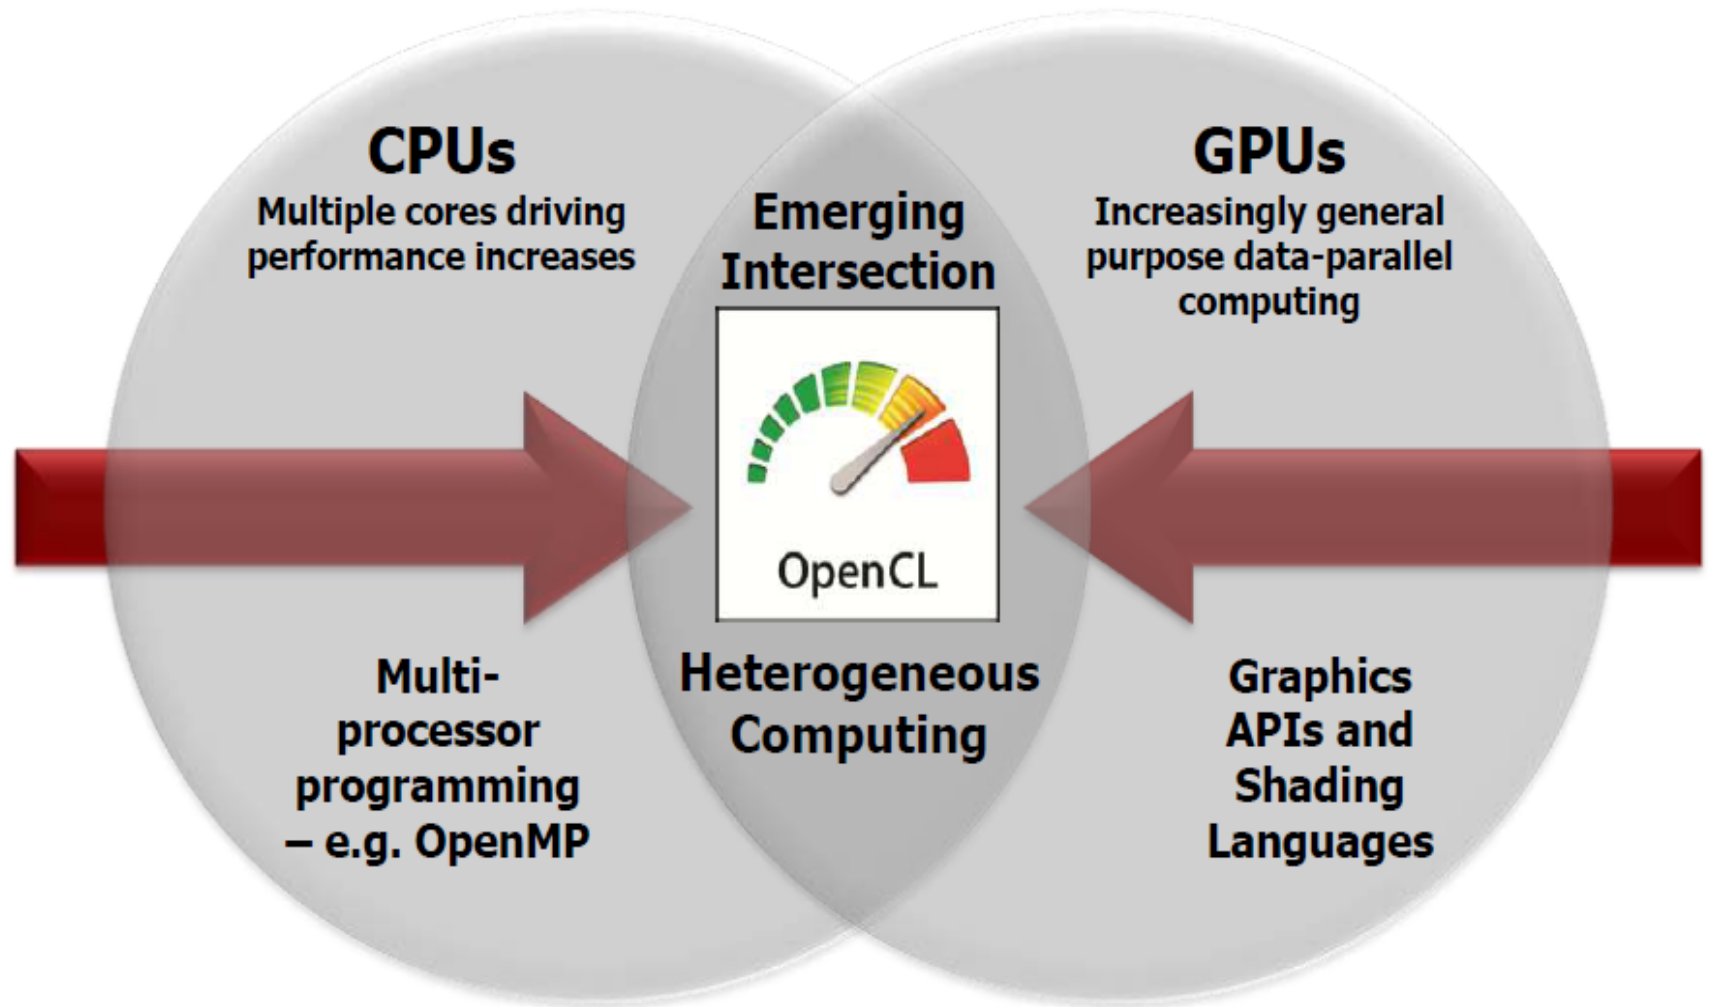
# Introduction and Overview

HPC Course

Excerpt of slides ©


**K H R O N O S**  
G R O U P<sup>TM</sup>

# Processor Parallelism



**OpenCL is a programming framework for heterogeneous compute resources**

# OpenCL Working Group



- **Diverse industry participation**
  - Processor vendors, system OEMs, middleware vendors, application developers
- **Many industry-leading experts involved in OpenCL's design**
  - A healthy diversity of industry perspectives
- **Apple made initial proposal and is very active in the working group**
  - Serving as specification editor



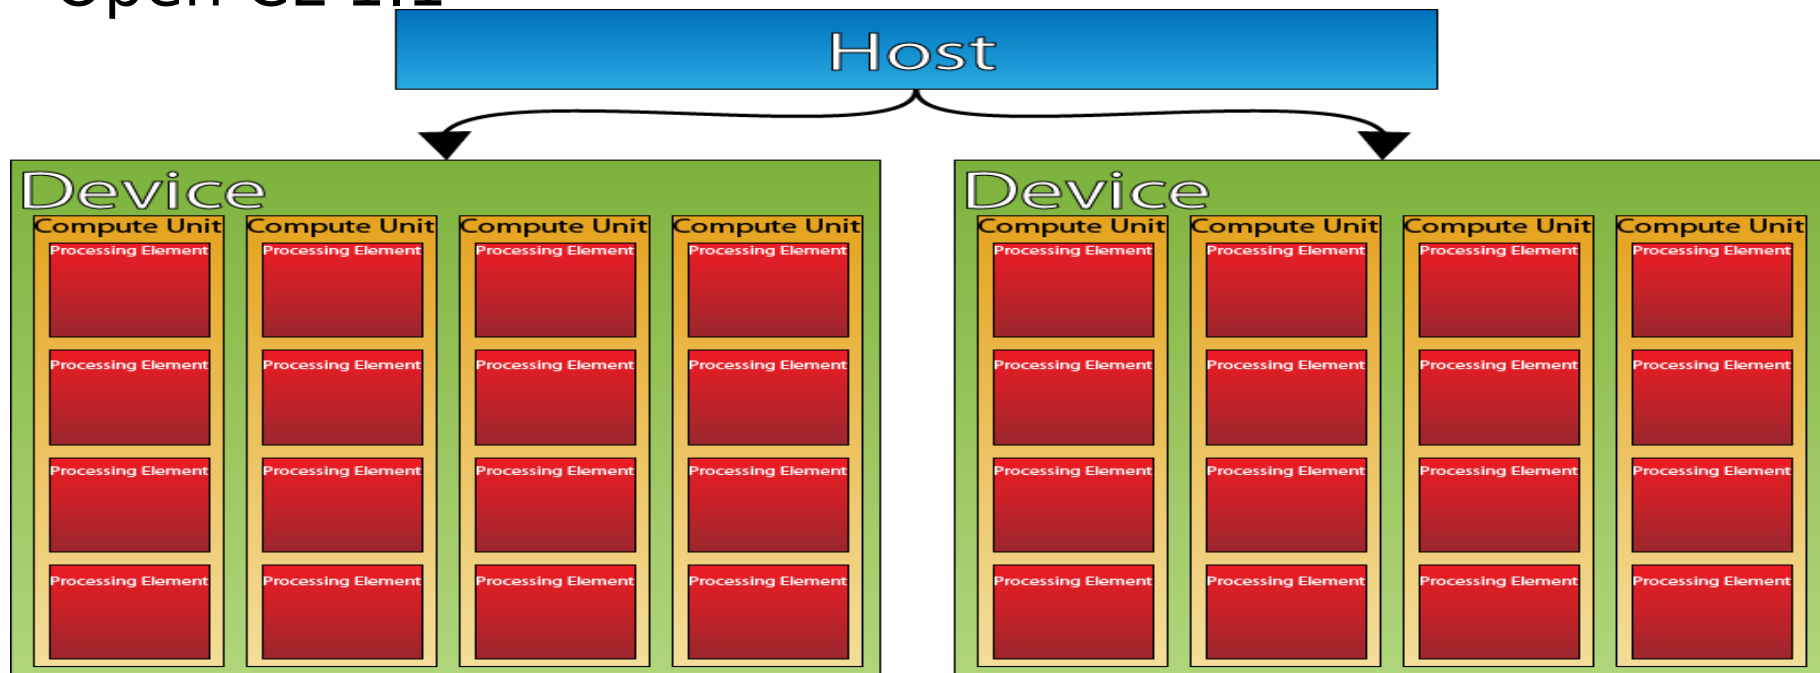



# OpenCL Timeline

2008 – Introduction of OpenCL by Apple

2008 – (Dec.) First specification of Open CL 1.0 by Khronos

2010 – Specification release and implementation ship  
Open CL 1.1

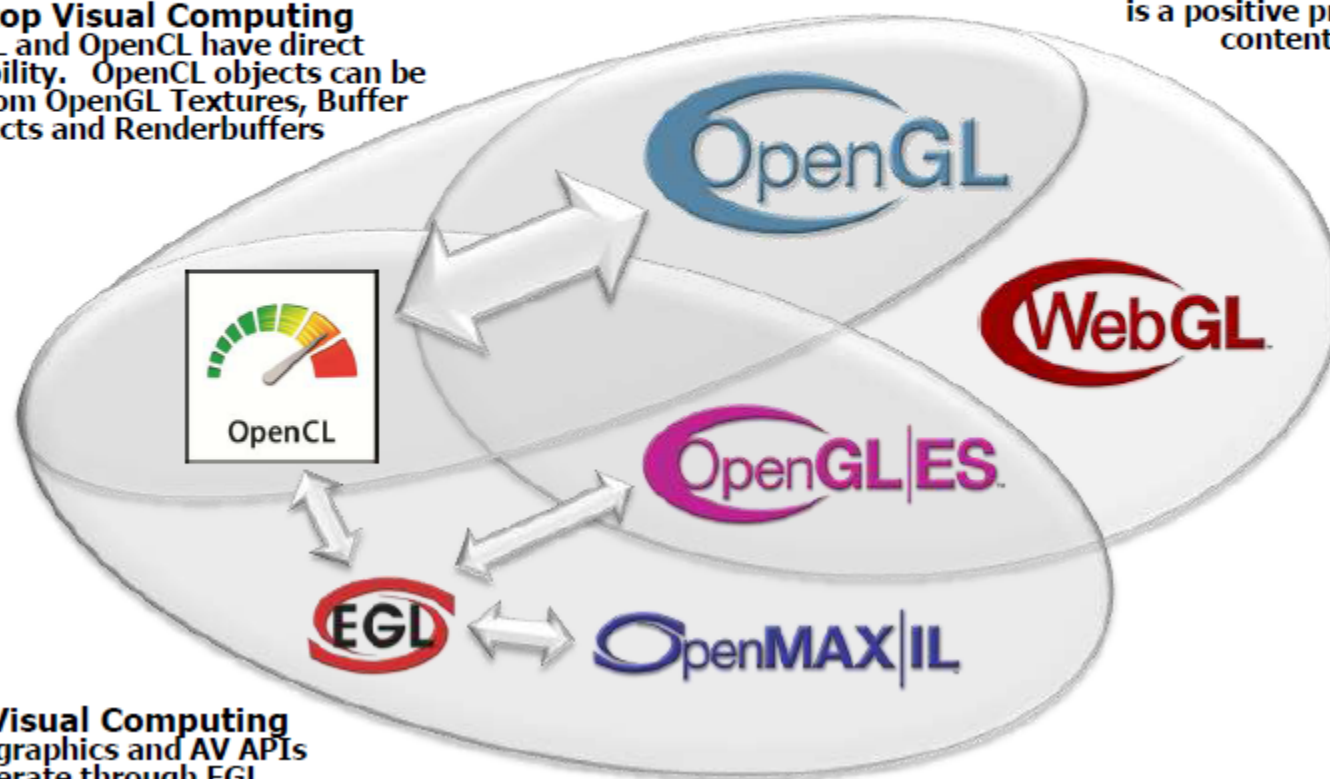




# OpenGL-based Ecosystem

**Desktop Visual Computing**  
OpenGL and OpenCL have direct interoperability. OpenCL objects can be created from OpenGL Textures, Buffer Objects and Renderbuffers

**Roadmap Convergence**  
OpenGL 4.0 and OpenGL ES 2.0 are both streamlined, programmable pipelines. GL and ES working groups are working on convergence. WebGL is a positive pressure for portable 3D content on all platforms



**Mobile Visual Computing**  
Compute, graphics and AV APIs interoperate through EGL

# OpenCL Overview

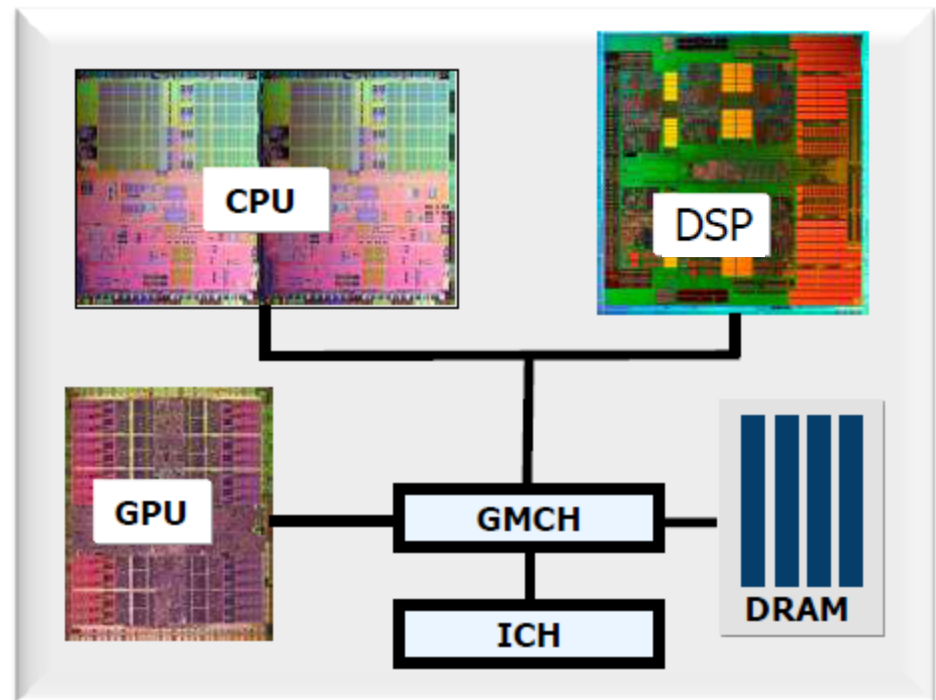


## It's a Heterogeneous World

- **A modern platform Includes:**

- One or more CPUs
- One or more GPUs
- DSP processors
- ... other?

OpenCL lets Programmers write a single portable program that uses ALL resources in the heterogeneous platform



GMCH = graphics memory control hub

ICH = Input/output control hub



# The BIG Idea behind OpenCL

- OpenCL execution model ...

**execute a kernel at each point in a problem domain**

- E.g., process a 1024 x 1024 image with one kernel invocation per pixel or  $1024 \times 1024 = 1,048,576$  kernel executions

## Traditional loops

```
void
trad_mul(int n,
        const float *a,
        const float *b,
        float *c)
{
    int i;
    for (i=0; i<n; i++)
        c[i] = a[i] * b[i];
}
```



## Data Parallel OpenCL

```
kernel void
dp_mul(global const float *a,
       global const float *b,
       global float *c)
{
    int id = get_global_id(0);



    c[id] = a[id] * b[id];
} // execute over "n" work-items
```





# An N-dimension domain of work-items

- **Define the "best" N-dimensioned index space for your algorithm**
  - Global Dimensions: 1024 x 1024 (whole problem space)
  - Local Dimensions: 128 x 128 (work group ... executes together)

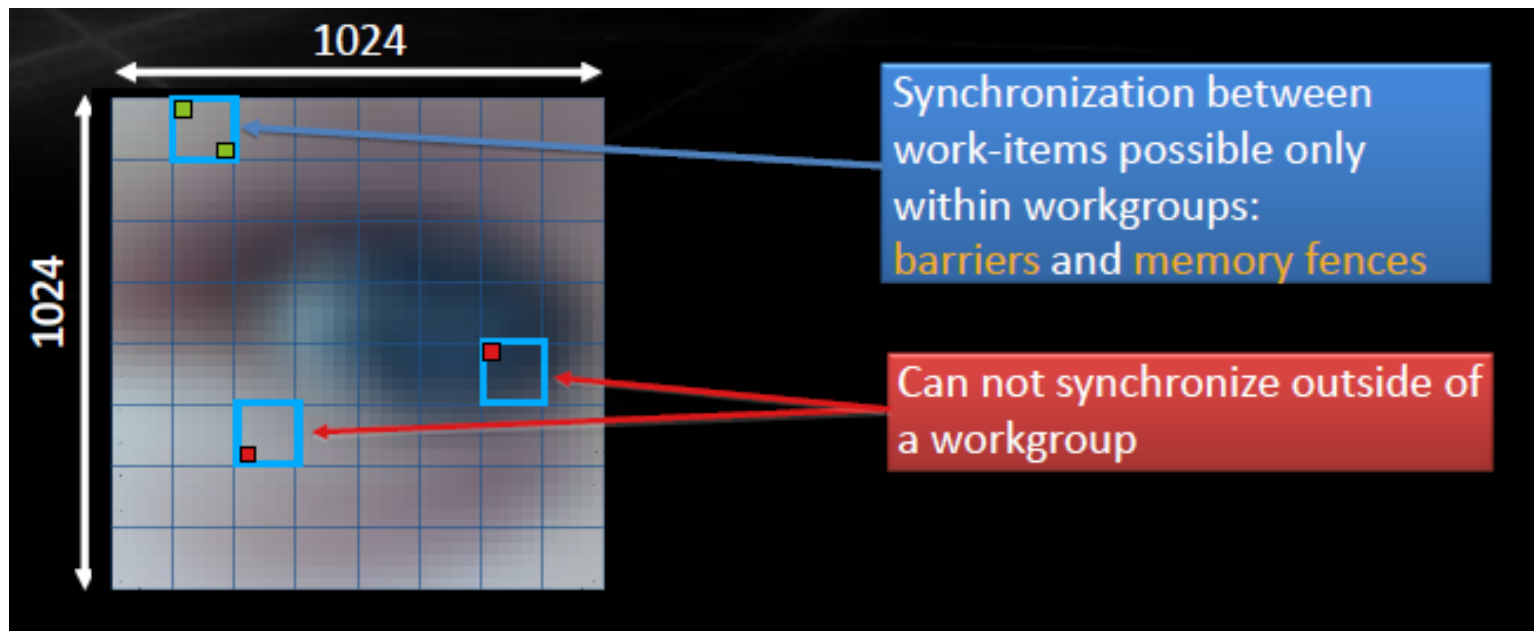






# An N-dimension domain of work-items

- **Define the “best” N-dimensioned index space for your algorithm**
  - Global Dimensions: 1024 x 1024 (whole problem space)
  - Local Dimensions: 128 x 128 (work group ... executes together)





# **To use OpenCL, you must**


- **Define the platform**
- **Execute code on the platform**
- **Move data around in memory**
- **Write (and build) programs**



# OpenCL Platform Model

- **One Host + one or more Compute Devices**


- Each Compute Device is composed of one or more Compute Units
- Each Compute Unit is further divided into one or more Processing Elements





# OpenCL Execution Model

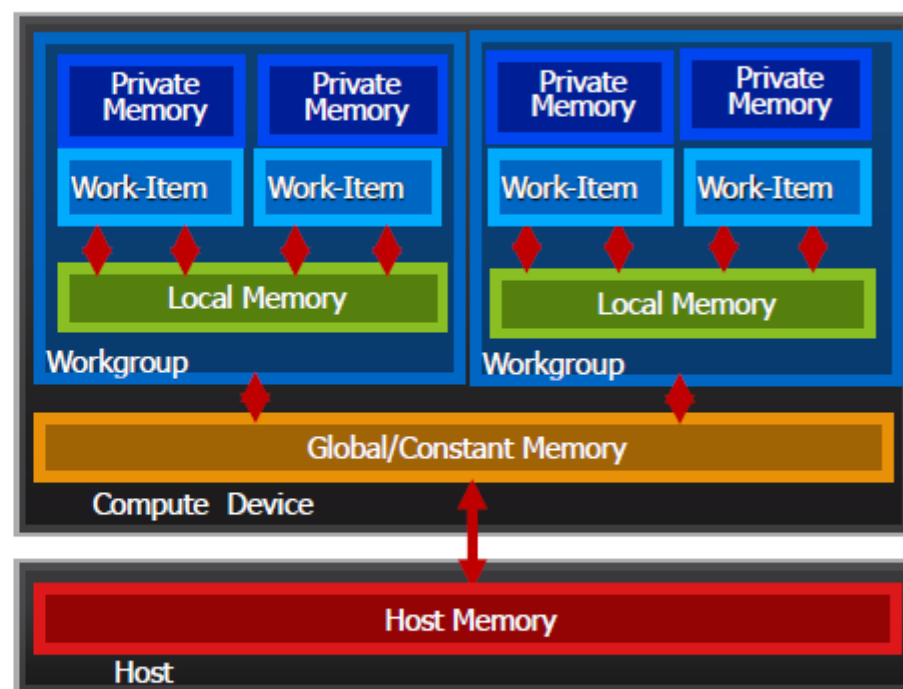
- **An OpenCL application runs on a host which submits work to the compute devices**
  - **Work item:** the basic unit of work on an OpenCL device
  - **Kernel:** the code for a work item. Basically a C function
  - **Program:** Collection of kernels and other functions (Analogous to a dynamic library)
  - **Context:** The environment within which work-items executes ... includes devices and their memories and command queues
- **Applications queue kernel execution instances**
  - Queued in-order ... one queue to a device
  - Executed in-order or out-of-order





# OpenCL Memory Model

- **Private Memory**
  - Per work-item
- **Local Memory**
  - Shared within a workgroup
- **Global/Constant Memory**
  - Visible to all workgroups
- **Host Memory**
  - On the CPU



**Memory management is Explicit**

You must move data from host -> global -> local ... *and* back



# Programming kernels: OpenCL C Language

- **A subset of ISO C99**
  - But without some C99 features such as standard C99 headers, function pointers, recursion, variable length arrays, and bit fields
- **A superset of ISO C99 with additions for:**
  - Work-items and workgroups
  - Vector types
  - Synchronization
  - Address space qualifiers
- **Also includes a large set of built-in functions**
  - Image manipulation
  - Work-item manipulation,
  - Specialized math routines, etc.



# Programming Kernels: Data Types

- **Scalar data types**

- char , uchar, short, ushort, int, uint, long, ulong, float
- bool, intptr\_t, ptrdiff\_t, size\_t, uintptr\_t, void, half (storage)

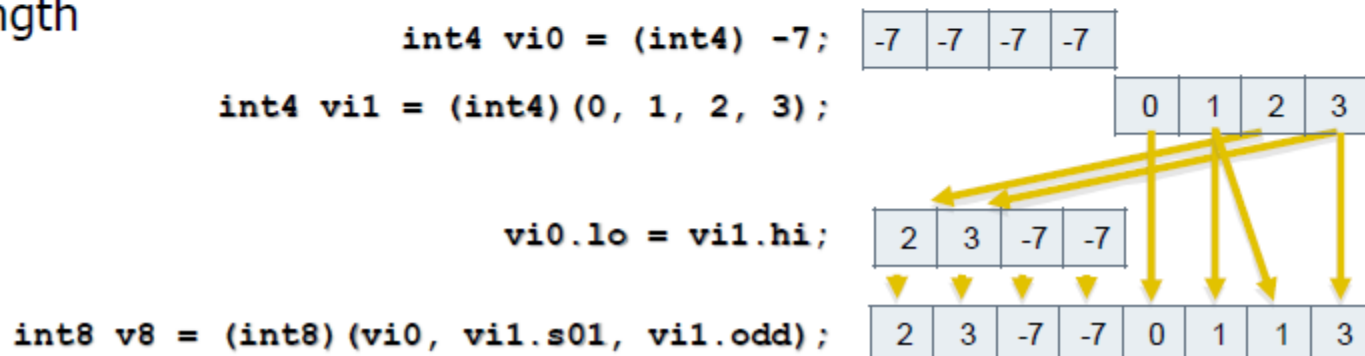
- **Image types**

- image2d\_t, image3d\_t, sampler\_t

- **Vector data types**

- Vector lengths 2, 4, 8, & 16 (char2, ushort4, int8, float16, double2, ...)
- Endian safe
- Aligned at vector length
- Vector operations
- Built-in functions

Double is an optional  
type in OpenCL 1.0







# Building Program Objects

- **The program object encapsulates:**
  - A context
  - The program source/binary
  - List of target devices and build options
- **The Build process ... to create a program object**
  - `clCreateProgramWithSource()`
  - `clCreateProgramWithBinary()`

## Kernel Code

```
kernel void
horizontal_reflect(read_only image2d_t src,
                  write_only image2d_t dst)
{
    int x = get_global_id(0); // x-coord
    int y = get_global_id(1); // y-coord
    int width = get_image_width(src);
    float4 src_val = read_imagef(src, sampler,
                                (int2)(width-1-x, y));
    write_imagef(dst, (int2)(x, y), src_val);
}
```



Compile for  
GPU



GPU  
code




Compile for  
CPU



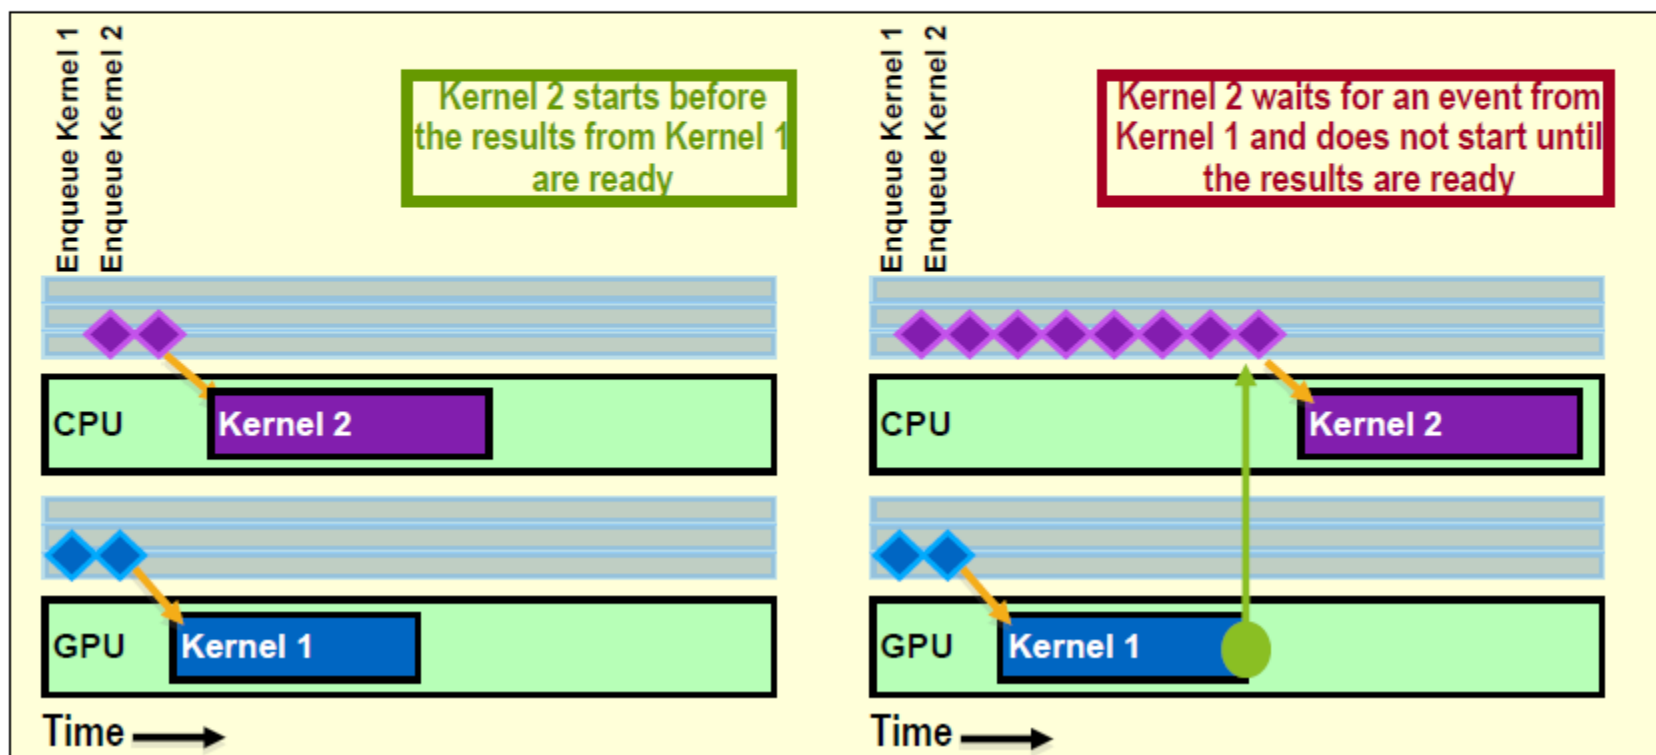
CPU  
code

Program



# OpenCL Synch: Queues & Events

- Events can be used to synchronize kernel executions between queues
- Example: 2 queues with 2 devices





# OpenCL Summary

